



Pattern Recognition

Lecture 15. Support Vector Machine

Dr. Shanshan ZHAO

School of AI and Advanced Computing
Xi'an Jiaotong-Liverpool University

Academic Year 2023-2024

Table of Contents

- ① SVM Intuition
- ② Formulization
- ③ Lagrange Duality
- ④ Kernel Trick
- ⑤ Soft Margin

Outline

- ① SVM Intuition
- ② Formulization
- ③ Lagrange Duality
- ④ Kernel Trick
- ⑤ Soft Margin

Introduction

- In the general case, the problem of finding linear discriminant functions can be formulated as a problem of **optimizing a criterion function**.
- Among all hyperplanes separating the data, there exists a **unique one yielding the maximum margin** of separation between the classes.

Binary Classification

Given training data (x_i, y_i) for $i = 1 \dots N$, with $x_i \in \mathbb{R}^d$ and $y_i \in \{-1, 1\}$, learn a classifier $f(x)$ such that



$$f(x_i) \begin{cases} \geq 0 & y_i = +1 \\ < 0 & y_i = -1 \end{cases}$$

i.e.



$y_i f(x_i) > 0 \rightarrow$ a correct classification

Linear separability

linearly
separable



not
linearly
separable




Linear Classifiers


A linear classifier has the form

$$f(x) = w^T x + b$$

- w is **normal**(vertical) to the line, and the b is the bias/intercept
 - *whether the positive of $f(x)$ is on the right or left of the line depends on the sign of the first parameter in vector w .*
- w is known as the **weight vector**.




(a) a line in 2D



(b) a plane in 3D


Linear Classifiers

- If training data is linearly separable, perceptron is guaranteed to find some linear separator/decision hyperplane.
- Which of these is optimal?



SVM Intuition

a very sensible choice for the hyperplane classifier would be the one that leaves the maximum margin from both classes.





Outline

- ① SVM Intuition
- ② Formulization
- ③ Lagrange Duality
- ④ Kernel Trick
- ⑤ Soft Margin

Margin

margin: a hyperplane leaves from both classes.

Our goal is to search for the direction that gives the maximum possible margin.

Recall that the distance of a point from a hyperplane is given by

$$z = \frac{|g(x)|}{||w||}$$

We can scale w, b so that the value of $g(x)$, at the nearest points in c_1, c_2 (circled in figure).

SVM objective

We can scale w, w_0 so that the value of $g(x)$, at the nearest points in c_1, c_2 (circled in figure1), is equal to 1 for class c_1 and equal to -1 for class c_2 , which is equivalent with

- 1. Having a margin of $\frac{1}{\|w\|} + \frac{1}{\|w\|} = \frac{2}{\|w\|}$
- 2. Requiring that

$$\begin{cases} w^T x + b \geq 1, & \forall x \in c_1 \\ w^T x + b \leq -1, & \forall x \in c_2 \end{cases}$$

- The support vectors lie on either of the two hyperplanes, that is

$$w^T x + b = \pm 1$$

Objective: Maximizing the margins

Maximize the margin (I) –Primal form(*)

- Optimization (Quadratic Programming) (known as a **Primal problem**.

$$\begin{cases} \text{minimize} & J(w, b) = \frac{1}{2} \|w\|^2 \\ \text{subject to} & y_i(w^T x_i + b) \geq 1, \quad i = 1, 2, \dots, N \end{cases} \quad (1)$$

- Minimizing the norm makes the margin maximum

It belongs to the convex programming family of problems, since the cost function is convex and the constraints are linear and define a convex set of feasible solutions. Such problems can be solved by considering the so-called Lagrangian duality.

Outline

- ① SVM Intuition
- ② Formulization
- ③ Lagrange Duality**
- ④ Kernel Trick
- ⑤ Soft Margin

Maximize the margin (II) –Dual form(*)

- The objective in Eq. (1) is a standard quadratic programming problem.
- Let $\boldsymbol{\lambda} \in \mathcal{R}^N$ be the dual variables, corresponding to Lagrange multipliers that enforce the N inequality constraints.

The generalized Lagrangian is given below

$$\mathcal{L}(\mathbf{w}, b, \boldsymbol{\lambda}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} - \sum_{i=1}^N \lambda_i [y_i (\mathbf{w}^T \mathbf{x}_i + b) - 1] \quad (2)$$

- where $\boldsymbol{\lambda}$ is the Lagrange multiplier, $\boldsymbol{\lambda} = (\lambda_1, \lambda_2, \dots, \lambda_N)^T$.
- we are **minimizing** with respect to \mathbf{w} and b , and **maximizing** with respect to $\boldsymbol{\lambda}$.

Dual problem *

We are **minimizing** with respect to \mathbf{w} and b , and **maximizing** with respect to λ .

The dual problem is

$$\max_{\lambda \geq 0} \min_{\mathbf{w}, b} \mathcal{L}(\mathbf{w}, b, \lambda)$$

Explanation:

- Appendix C.4 from *Pattern Recognition by Segios*
- cs229-notes3 by *Andrew Ng*, page 7-9.

<https://core.xjtlu.edu.cn/mod/folder/view.php?id=69950>

Maximize the margin (II) –Dual form(*)

Lets find the dual form of the problem. To do so, we need to first minimize $L(\mathbf{w}, b, \lambda)$ with respect to w and b (for fixed λ), which we'll do by setting the derivatives of L with respect to \mathbf{w} and b to zero. We obtain the following two conditions

$$\frac{\partial}{\partial \mathbf{w}} \mathcal{L}(\mathbf{w}, b, \lambda) = 0 \quad (3)$$

$$\frac{\partial}{\partial b} \mathcal{L}(\mathbf{w}, b, \lambda) = 0 \quad (4)$$

(5)

Combining (3) (4) and (2), results in

$$\mathbf{w} = \sum_{i=1}^N \lambda_i y_i \mathbf{x}_i \quad (6)$$

$$\sum_{i=1}^N \lambda_i y_i = 0 \quad (7)$$

Support Vectors

The Lagrange multipliers can be either zero or positive (Appendix C). Thus, the vector parameter \mathbf{w} of the optimal solution is a linear combination of $N_s \leq N$ feature vectors that are associated with $\lambda_i \neq 0$. That is,

$$\mathbf{w} = \sum_{i=1}^{N_s} \lambda_i y_i \mathbf{x}_i$$

These are known as **support vectors** and the optimum hyperplane classifier as a support vector machine (SVM). As it is pointed out in Appendix C, a nonzero Lagrange multiplier corresponds to a so called active constraint.

Hence, as the set of constraints in equation (7) suggests for $\lambda_i \neq 0$, the support vectors lie on either of the two hyperplanes, that is,

$$\mathbf{w}^T \mathbf{x} + b = \pm 1$$

Maximize the margin (II) –Dual form(*)

Plugging equation (6) and (7) into Lagrangian Equation (2) yields the following

$$\begin{aligned}\mathcal{L}(w, b, \lambda) &= \frac{1}{2}w^T w - \sum_{i=1}^N \lambda_i y_i w^T x_i - \sum_{i=1}^N \lambda_i y_i b + \sum_{i=1}^N \lambda_i \\ &= \frac{1}{2}w^T w - w^T w - 0 + \sum_{i=1}^N \lambda_i \\ &= -\frac{1}{2}w^T w + \sum_{i=1}^N \lambda_i \\ &= -\frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \lambda_i \lambda_j y_i y_j x_i^T x_j + \sum_{n=1}^N \lambda_i\end{aligned}$$

Recall that we got to the equation above by minimizing L with respect to \mathbf{w} and b . Putting this together with the constraints $\lambda_i \geq 0$ and the constraint $\sum_{i=1}^N \lambda_i y_i = 0$, we obtain the following dual optimization problem:

$$\max_{\lambda} W(\lambda) = \max_{\lambda} \left(\sum_{i=1}^N \lambda_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \lambda_i \lambda_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j \right) \quad (8)$$

$$\text{subject to } \sum_{i=1}^N \lambda_i y_i = 0 \quad (9)$$

$$\lambda \geq 0 \quad (10)$$

We should be able to verify that the Karush-Kuhn-Tucker (KKT) conditions to hold are indeed satisfied in our optimization problem.

- Hence, we can solve the dual in lieu of solving the primal problem. Specifically, in the dual problem above, we have a maximization problem in which the parameters are the λ_i .
- the specific algorithm that we're going to use to solve the dual problem is : **SMO**.
- Most materials leave this part (how to find λ) to the end, but turns to introduce the kernels. *Don't get lost*.
- Details of *SMO* algorithm: cs229-notes3 by Andrew Ng, page 20-25.
- Let's leave it alone, as well :)

- If we are indeed able to solve it (i.e., find the λ_i that maximize $W(\lambda)$ subject to the constraints)
- then we can use Equation (6) $\mathbf{w} = \sum_{i=1}^N \lambda_i y_i \mathbf{x}_i$ to go back and find the optimal \mathbf{w} as a function of the λ .
- Having found \mathbf{w}^* , by considering the primal problem, it is also straightforward to find the optimal value for the intercept term b .
- In practice, b is computed as an average value obtained using all conditions of this type.

SVM


In practice, b is computed as an average value obtained using all conditions of this type.

$$\begin{aligned} b &= \frac{1}{|\mathcal{S}|} \sum_{i \in \mathcal{S}} (y_i - \mathbf{w}^T \mathbf{x}_i) \\ &= \frac{1}{|\mathcal{S}|} \sum_{i \in \mathcal{S}} (y_i - \sum_{j \in \mathcal{S}} \lambda_j y_j \mathbf{x}_j^T \mathbf{x}_i) \\ &= \frac{1}{|\mathcal{S}|} \sum_{i \in \mathcal{S}} (y_i - \sum_{j \in \mathcal{S}} \lambda_j y_j \langle \mathbf{x}_j, \mathbf{x}_i \rangle) \end{aligned}$$

where \mathcal{S} is the set of support vectors. We end up with the equivalent optimization task.

SVM

linearly separable data



Outline

- ① SVM Intuition
- ② Formulization
- ③ Lagrange Duality
- ④ Kernel Trick**
- ⑤ Soft Margin

- Suppose we've fit our model's parameters to a training set, and now wish to **make a prediction at a new point input** x .
- We would then calculate $w^T x + b$, and predict $y = 1$ if and only if this quantity is bigger than zero. With equation (6), this quantity can also be written:

$$w^T x + b = \left(\sum_{i=1}^{N_s} \lambda_i y_i x_i \right)^T x + b \quad (11)$$


$$= \sum_{i=1}^{N_s} \lambda_i y_i \langle x_i, x \rangle + b \quad (12)$$

Hence, if we've found the λ_i , in order to make a prediction, we have to calculate a quantity that depends only on the **inner product** between x and the points in the training set.

- Moreover, we saw earlier that the λ_i will all be zero except for the support vectors.
- Thus, many of the terms in the sum above will be zero, and we really need to find only the inner products between x and the support vectors (of which there is often only a small number) in order calculate equation (12) and make our prediction.
- What if we have non-linear cases in the original space?
 - We discussed that we can map x into higher space where exists a linear hyperplane.

Kernel trick : Feature mapping

- Rather than applying SVMs using the original input attributes x , we may instead want to learn using some features $\phi(x)$.



SVM : Kernel trick

- Rather than applying SVMs using the original input attributes x , we may instead want to learn using some features $\phi(x)$.
- To do so, we simply need to go over our previous algorithm, and replace x everywhere in it with $\phi(x)$.
- Since the algorithm can be written entirely in terms of the inner products $\langle x, z \rangle$, this means that we would replace all those inner products with $\langle \phi(x), \phi(z) \rangle$.

SVM : Kernel trick

- Both the quadratic programming problem and the final decision function

$$g(x) = \text{sign}\left(\sum_{i=1}^n \lambda_i y_i \langle x \cdot x_i \rangle + b\right) \quad (13)$$

depend only on the dot products.

- We can generalize this result to the non-linear case by mapping the original input space into some other space \mathcal{F} using a non-linear map $\phi : \mathcal{R}^d \rightarrow \mathcal{F}$ and perform the linear algorithm in the \mathcal{F} space which only requires the inner products.

$$k(x, y) = \phi(x)^T \phi(y)$$

SVM : Kernel trick

- This results in the non-linear decision function of the form

$$g(x) = \text{sign}\left(\sum_{i=1}^n \lambda_i y_i k(x, x_i) + b\right) \quad (14)$$

where the parameters λ_i are computed as the solution of the quadratic programming problem.

- In the original input space, the hyperplane corresponds to a non-linear decision function whose form is determined by the kernel.
- We can use $k(x, x')$ directly for computation without transforming x and x' , as long as ϕ exists.**

Example 1:

Suppose $x, z \in \mathcal{R}^2$, i.e., $x = [x_1, x_2]^T$, $z = [z_1, z_2]^T$,
consider

$$k(x, z) = (x^T z)^2$$

$k(x, z)$ would be a valid kernel if we can find a projection function $\phi(x)$ that satisfy $k(x, z) = \phi(x)^T \phi(z)$

We may check $\phi(x) = [x_1^2, \sqrt{2}x_1x_2, x_2^2]^T$ (yes, this is kind of cheating, but we just want to understand how kernel function works for now.)

$$\begin{aligned}\phi(a)^T \phi(b) &= [a_1^2, \sqrt{2}a_1a_2, a_2^2]^T [b_1^2, \sqrt{2}b_1b_2, b_2^2] \\ &= a_1^2b_1^2 + 2a_1a_2b_1b_2 + a_2^2b_2^2\end{aligned}$$

$$\begin{aligned}k(a, b) &= (a^T b)^2 \\ &= ([a_1, a_2]^T [b_1, b_2])^2 \\ &= (a_1b_1 + a_2b_2)^2 \\ &= a_1^2b_1^2 + 2a_1a_2b_1b_2 + a_2^2b_2^2\end{aligned}$$

So,

$$k(a, b) = \phi(a)^T \phi(b)$$

Example 2.

what if $x, z \in \mathcal{R}^n$?

page 14 on *cs229-notes3* by Andrew Ng

Kernel : inner product

Polynomial kernels

$$\forall x, x' \in \mathcal{R}^2,$$

$$\begin{aligned} k(x, x') &= (1 + x^T x')^2 = (1 + x_1 x'_1 + x_2 x'_2)^2 \\ &= 1 + x_1^2 x'^2_1 + x_2^2 x'^2_2 + 2x_1 x'_1 + 2x_2 x'_2 + 2x_1 x'_1 x_2 x'_2 \end{aligned}$$

↓ above is an inner product of two vectors

$$\begin{aligned} &= \begin{bmatrix} 1 \\ x_1^2 \\ x_2^2 \\ \sqrt{2}x_1 \\ \sqrt{2}x_2 \\ \sqrt{2}x_1x_2 \end{bmatrix} \cdot \begin{bmatrix} 1 \\ x'^2_1 \\ x'^2_2 \\ \sqrt{2}x'_1 \\ \sqrt{2}x'_2 \\ \sqrt{2}x'_1x'_2 \end{bmatrix} \\ &= z^T z' \end{aligned}$$

with $z = \phi(x) = [1, x_1^2, x_2^2, \sqrt{2}x_1, \sqrt{2}x_2, \sqrt{2}x_1x_2]$

Example




Figure 5.2 Illustration of the XOR classification problem and the use of polynomial kernels. (a) XOR problem linearly non-separable in the input space. (b) Linearly separable using second-degree polynomial kernel.

SVM : Kernel trick

- Even though \mathcal{F} may be high-dimensional, a simple kernel $k(x, y)$ such as the following can be computed efficiently.

Polynomial

$$k(\mathbf{x}, \mathbf{y}) = (\mathbf{x} \cdot \mathbf{y})^p$$

Sigmoidal

$$k(\mathbf{x}, \mathbf{y}) = \tanh(\kappa(\mathbf{x} \cdot \mathbf{y}) + \theta)$$

Radial basis function

$$k(\mathbf{x}, \mathbf{y}) = \exp(-\|\mathbf{x} - \mathbf{y}\|^2 / (2\sigma^2))$$

Figure 2: Common kernel functions

- Once a kernel function is chosen, we can substitute $\phi(x_i)$ for each training example x_i , and perform the optimal hyperplane algorithm in \mathcal{F} .

Kernel Matrix

consider some finite set of m points (not necessarily the training set) $\{x(1), \dots, x(m)\}$, and let a square, m -by- m matrix K be defined so that its (i, j) -entry is given by $K_{ij} = K(x(i), x(j))$. This matrix is called the **Kernel matrix**.

- K must be symmetric
- K is positive semi-definite.

It turns out to be a necessary and a sufficient condition for k to be a valid kernel (also called a Mercer kernel) if its corresponding Kernel Matrix is symmetric positive semidefinite.

- proof also in the Andrew Ng's note.

Theorem(Mercer)

Let $k : \mathcal{R}^n \times \mathcal{R}^n \mapsto \mathcal{R}$ be given. Then for k to be a valid (Mercer) kernel, it is necessary and sufficient that for any $\{x(1), \dots, x(m)\}$, ($m < \infty$), the corresponding kernel matrix is symmetric positive semi-definite.

Given a function k , apart from trying to find a feature mapping ϕ that corresponds to it, this theorem therefore gives another way of testing if it is a valid kernel.



Outline

- ① SVM Intuition
- ② Formulization
- ③ Lagrange Duality
- ④ Kernel Trick
- ⑤ Soft Margin

Regularization and the non-separable case

- The derivation of the SVM as presented so far assumed that the data is linearly separable. While mapping data to a high dimensional feature space via ϕ does generally increase the likelihood that the data is separable, we can't guarantee that it always will be so.
- Also, in some cases it is not clear that finding a separating hyperplane is exactly what we'd want to do, since that might be susceptible to outliers.

Regularization and the non-separable case



the left figure below shows an optimal margin classifier, and when a single outlier is added in the upper-left region (right figure), it causes the decision boundary to make a dramatic swing, and the resulting classifier has a much smaller margin.

The training feature vectors now belong to one of the following three categories:

- 1 Vectors that fall outside the band and are correctly classified. These vectors comply with the constraints

$$y_i(w^T x_i + b) \geq 1, i = 1, 2, \dots, N$$

- 2 Vectors falling inside the band and are correctly classified. These are the points placed in circles in Figure 3, and they satisfy the inequality

$$0 \leq y_i(w^T x_i + b) < 1$$

- 3 Vectors that are misclassified. They are enclosed by circles and obey the inequality

$$y_i(w^T x_i + b) < 0$$


All three cases can be treated under a single type of constraints by introducing a new set of variables, namely,

$$y_i(w^T x_i + b) \geq 1 - \xi_i$$

slack variables

$$y_i(w^T x_i + b) \geq 1 - \xi_i$$

The first category of data corresponds to $\xi = 0$, the second to $0 < \xi_i \leq 1$, and the third to $\xi_i > 1$. The variables ξ_i are known as **slack variables**. The slack variables ξ_i are used to handle misclassified instances.



Regularization and the non-separable case

The goal now is to make the margin as large as possible but at the same time to keep the number of points with $\xi_i > 0$ as small as possible.

we reformulate our optimization

$$\begin{cases} \text{minimize} & J(w, b, \xi) = \frac{1}{2} \|w\|^2 + C \sum_{i=1}^N \xi_i \\ \text{subject to} & y_i(w^T x_i + b) \geq 1 - \xi_i, \quad i = 1, 2, \dots, N \\ & \xi_i \geq 0, i = 1, \dots, N \end{cases} \quad (15)$$

- This is still a quadratic optimization problem and there is a unique minimum.
- Thus, samples are now permitted to have (functional) margin less than 1, and if a sample whose functional margin is $1 - \xi_i$, we would pay a cost of the objective function being increased by $C\xi_i$.

Regularization parameter C

- The term $C \sum_{i=1}^n \xi_i$ can be thought of as measuring some amount of misclassification where lowering the value of C corresponds to a smaller penalty for misclassification.
- The parameter $C > 0$ controls the trade-off between the slack variable penalty and the margin.
- Because any point that is misclassified has $\xi_i > 1$, it follows that $\sum_i \xi_i$ is an upper bound on the number of misclassified points.
- The parameter C is therefore analogous to a regularization coefficient because it controls the trade-off between minimizing training errors and controlling model complexity.
- In the limit $C \rightarrow \infty$, we will recover the earlier support vector machine for separable data.

we can form the Lagrangian:

$$\mathcal{L}(\mathbf{w}, b, \xi, \lambda, \gamma) = \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^N \xi_i - \sum_{i=1}^N \lambda_i [y_i (\mathbf{w}^T \mathbf{x} + b) - 1 + \xi_i] - \sum_{i=1}^N \gamma_i \xi_i$$

the λ_i and γ_i are our Lagrange multipliers. (constrained to be ≥ 0).

We won't go through the derivation of the dual again in detail, but after setting the derivatives with respect to \mathbf{w} and b to zero as before, substituting them back in, and simplifying,

we obtain the following dual form of the problem:

$$\max_{\lambda} W(\lambda) = \max_{\lambda} \left(\sum_{i=1}^N \lambda_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \lambda_i \lambda_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j \right) \quad (16)$$

$$\text{subject to } 0 \leq \lambda_i \leq C, i = 1, \dots, N \quad (17)$$

$$\sum_{i=1}^N \lambda y_i = 0 \quad (18)$$

The corresponding set of KKT conditions can be given and used for testing for the convergence of the SMO algorithm.

SVM Training Methodology

- Training is formulated as an optimization problem
 - Dual problem - Makes use of Lagrange multipliers
 - Kernel trick
- Determination of the model parameters corresponds to a convex optimization problem
 - Solution is straightforward (local solution is a global optimum)
- Noisy labels – Soft Margin

Check the code.

Requirements

You should be able to describe

- the idea behind SVM.
- concept of the kernel trick, use examples.
- how the regularization parameter C works in SVM.

Thank You !
Q & A